SIGMXD/ PXDS AMSTERDAM 2X19

Chair: Norman May

A Scalable Index for Top-k Subtree Similarity Queries
Daniel Kocher (University of Salzburg), Nikolaus Augsten (University of Salzburg)

A Layered Aggregate Engine for Analytics Workloads
Maximilian Schleich (University of Oxford), Dan Olteanu (University of Oxford), Mahmoud Abo Khamis (RelationalAI), Hung Ngo (RelationalAI), XuanLong Nguyen (University of Michigan)

Towards Scalable Hybrid Stores: Constraint-Based Rewriting to the Rescue
Rana Alotaibi (University of California, San Diego), Damian Bursztyn (Thales), Alin Deutsch (University of California, San Diego), Ioana Manolescu (Inria & Ecole polytechnique), Stamatis Zampetakis (Orchestra Networks)

MIFO: A Query-Semantic Aware Resource Allocation Policy
Prajakta Kalmegh (Duke University), Shivnath Babu (Unravel Data Systems)

Dissecting the Performance of Strongly-Consistent Replication Protocols
Ailidani Ailijiang (Microsoft), Aleksey Charapko (University at Buffalo, SUNY), Murat Demirbas (University at Buffalo, SUNY)

FishStore: Faster Ingestion with Subset Hashing
Dong Xie (University of Utah), Badrish Chandramouli (Microsoft Research), Yinan Li (Microsoft Research), Donald Kossmann (Microsoft Research)