

# SIGMOD/

# AMSTERDAM

# PXDS

# 2X19

**Thursday 14:20-16:20 – Administratiezaal 1.1**

**SIGMOD Research 17: Scalability**

**Chair:** Norman May

**A Scalable Index for Top-k Subtree Similarity Queries**

*Daniel Kocher (University of Salzburg), Nikolaus Augsten (University of Salzburg)*

**A Layered Aggregate Engine for Analytics Workloads**

*Maximilian Schleich (University of Oxford), Dan Olteanu (University of Oxford), Mahmoud Abo Khamis (RelationalAI), Hung Ngo (RelationalAI), XuanLong Nguyen (University of Michigan)*

**Towards Scalable Hybrid Stores: Constraint-Based Rewriting to the Rescue**

*Rana Alotaibi (University of California, San Diego), Damian Bursztyn (Thales), Alin Deutsch (University of California, San Diego), Ioana Manolescu (Inria & Ecole polytechnique), Stamatis Zampetakis (Orchestra Networks)*

**MIFO: A Query-Semantic Aware Resource Allocation Policy**

*Prajakta Kalmegh (Duke University), Shivnath Babu (Unravel Data Systems)*

**Dissecting the Performance of Strongly-Consistent Replication Protocols**

*Ailidani Ailijiang (Microsoft), Aleksey Charapko (University at Buffalo, SUNY), Murat Demirbas (University at Buffalo, SUNY)*

**FishStore: Faster Ingestion with Subset Hashing**

*Dong Xie (University of Utah), Badrish Chandramouli (Microsoft Research), Yinan Li (Microsoft Research), Donald Kossmann (Microsoft Research)*